

# Gradient Boosted Decision Trees (GBDT)

## Tutorial on Implementation and Interpretability in MoDeVa

July 5, 2025

### **Tutorial Overview:**

- GBDT fundamentals and mathematical foundations
- Implementation with XGBoost, LightGBM, and CatBoost
- Interpretability through Functional ANOVA
- Special cases: GAM and GAMI structures
- Monotonicity constraints and domain knowledge
- Global and local interpretation techniques

# Outline

- 1 Introduction to GBDT
- 2 Mathematical Foundations
- 3 Implementation in MoDeVa
- 4 Interpretability Through Functional ANOVA
- 5 Global Interpretation Methods
- 6 Local Interpretation Methods
- 7 Monotonicity Constraints
- 8 Best Practices and Guidelines
- 9 Conclusion

# What is Gradient Boosted Decision Trees?

**Definition:** GBDT is a powerful ensemble learning algorithm that builds a sequence of decision trees, where each subsequent tree is trained to correct the errors of its predecessors.

## Key Characteristics:

- **Sequential Learning:** Trees built one after another
- **Error Correction:** Each tree learns from previous mistakes
- **Gradient-Based:** Uses gradients of loss function for optimization
- **Versatile:** Handles both regression and classification tasks

## vs. Random Forests:

- GBDT: Sequential, dependent trees
- Random Forests: Parallel, independent trees

# Why GBDT Matters

## Advantages:

- **High Accuracy:** Often achieves state-of-the-art performance
- **Flexibility:** Works well with various data types
- **Feature Handling:** Robust to missing values and outliers
- **Interpretability:** Can be made interpretable with constraints

## Applications:

- Credit scoring and risk assessment
- Sales forecasting and demand prediction
- Medical diagnosis and treatment planning
- Marketing response modeling

## Popular Implementations:

- XGBoost: Performance and speed optimization
- LightGBM: Memory efficiency and fast training
- CatBoost: Categorical feature handling

## Overall Ensemble Model:

$$F_M(x) = F_0(x) + \sum_{m=1}^M \gamma_m T_m(x) \quad (1)$$

## Where:

- $F_0(x)$ : Initial model (often constant, e.g., mean for regression)
- $T_m(x)$ : Decision tree added at iteration  $m$
- $\gamma_m$ : Learning rate controlling tree contribution
- $M$ : Total number of trees

**Key Insight:** Each tree adds a small correction to the ensemble, gradually improving predictions through iterative refinement.

# Pseudo-Residual Calculation

## Gradient-Based Learning:

At each iteration  $m$ , compute pseudo-residuals:

$$r_{im} = -\frac{\partial L(y_i, F_{m-1}(x_i))}{\partial F_{m-1}(x_i)} \quad (2)$$

## Components:

- $y_i$ : True target value for sample  $i$
- $F_{m-1}(x_i)$ : Current model prediction for sample  $i$
- $r_{im}$ : Pseudo-residual (negative gradient) for sample  $i$
- $L$ : Loss function (e.g., squared loss, log-loss)

**Intuition:** Pseudo-residuals point in the direction of steepest descent, indicating how to adjust predictions to reduce loss.

# Model Update Rule

## Sequential Update Process:

After fitting tree  $h_m(x)$  to pseudo-residuals:

$$F_m(x) = F_{m-1}(x) + \gamma_m h_m(x) \quad (3)$$

## Complete GBDT Algorithm:

- ① Initialize:  $F_0(x) = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, \gamma)$
- ② For  $m = 1$  to  $M$ :
  - Compute pseudo-residuals:  $r_{im} = -\frac{\partial L(y_i, F_{m-1}(x_i))}{\partial F_{m-1}(x_i)}$
  - Fit tree  $h_m(x)$  to residuals
  - Find optimal step size:  $\gamma_m = \arg \min_{\gamma} \sum_{i=1}^n L(y_i, F_{m-1}(x_i) + \gamma h_m(x_i))$
  - Update model:  $F_m(x) = F_{m-1}(x) + \gamma_m h_m(x)$
- ③ Return final model:  $F_M(x)$

## Unified Wrapper Architecture:

MoDeVa provides comprehensive wrappers around leading GBDT implementations:

### XGBoost

- Speed optimization
- Advanced regularization
- Cross-validation built-in
- Monotonicity constraints

### LightGBM

- Memory efficient
- Fast training speed
- Categorical features
- Large-scale datasets

### CatBoost

- Categorical handling
- Overfitting resistant
- No preprocessing
- Symmetric trees

## Benefits of Unified Interface:

- Consistent API across implementations
- Easy model comparison and switching
- Integrated interpretability tools
- Unified hyperparameter management



# Data Setup and Preprocessing

## Preparing Data for GBDT:

```
1 from modeva import DataSet
2
3 # Create dataset object
4 ds = DataSet()
5 # Load built-in dataset
6 ds.load(name="BikeSharing")
7
8 # Preprocessing pipeline
9 ds.scale_numerical(features=("cnt",), method="log1p")
10 ds.set_feature_type(feature="hr",
11                     feature_type="categorical")
12 ds.set_feature_type(feature="mnth",
13                     feature_type="categorical")
14 ds.scale_numerical(features=ds.feature_names_numerical,
15                     method="standardize")
16 ds.set_inactive_features(features=("yr", "season", "temp"))
17 ds.preprocess()
18
19 # Split data
20 ds.set_random_split(test_ratio=0.2, random_state=42)
```

# Model Configuration

## Setting Up GBDT Models: Regression Models:

```
1 from modeva.models import (MoLGBMRegressor,
2                             MoXGBRegressor,
3                             MoCatBoostRegressor)
4
5 # LightGBM
6 model_lgbm = MoLGBMRegressor(name="LGBM_model",
7                               max_depth=2,
8                               n_estimators=100)
9
10 # XGBoost
11 model_xgb = MoXGBRegressor(name="XGB_model",
12                             max_depth=2,
13                             n_estimators=100)
14
15 # CatBoost
16 model_cat = MoCatBoostRegressor(name="CBoost_model",
17                                  max_depth=2,
18                                  n_estimators=100)
```

## Classification Models:

```
1 from modeva.models import (MoLGBMClassifier,  
2                             MoXGBClassifier,  
3                             MoCatBoostClassifier)  
4  
5 # Similar syntax for classification tasks  
6 model_clf = MoXGBClassifier(name="XGB_classifier",  
7                             max_depth=2,  
8                             n_estimators=100)
```

# Model Training and Evaluation

## Training Process:

```
1 # Train the model
2 model_gbdt.fit(ds.train_x, ds.train_y)
3
4 # Make predictions
5 train_pred = model_gbdt.predict(ds.train_x)
6 test_pred = model_gbdt.predict(ds.test_x)
```

## Performance Assessment:

```
1 from modeva import TestSuite
2 # Create comprehensive test suite
3 ts = TestSuite(ds, model_gbdt)
4 # Performance metrics
5 result = ts.diagnose_accuracy_table()
6 print(result.table)
```

## Performance Metrics Include:

- Regression: RMSE, MAE,  $R^2$ , MAPE
- Classification: Accuracy, Precision, Recall, F1, AUC
- Cross-validation scores
- Training vs. validation performance

# Making GBDT Interpretable

**The Challenge:** GBDT models are typically considered "black boxes"

**The Solution:** Functional ANOVA Decomposition

**Mathematical Framework:**

$$f(x) = \mu + \sum_j f_j(x_j) + \sum_{j < k} f_{jk}(x_j, x_k) + \dots \quad (4)$$

**Components:**

- $\mu$ : Global intercept (average prediction)
- $f_j(x_j)$ : Main effects (individual feature contributions)
- $f_{jk}(x_j, x_k)$ : Pairwise interaction effects
- Higher-order terms: Complex interactions (limited by tree depth)

**Key Insight:** Deep trees create complex interactions; shallow trees enable clearer interpretation.

# Special Cases: GAM and GAMI Structures

## Depth-1 Trees: GAM Structure

$$f(x) = \mu + \sum_j f_j(x_j) \quad (5)$$

- Each tree makes single split: Results in Generalized Additive Model
- Only main effects, no interactions
- Highly interpretable but limited flexibility

## Depth-2 Trees: GAMI Structure

$$f(x) = \mu + \sum_j f_j(x_j) + \sum_{j < k} f_{jk}(x_j, x_k) \quad (6)$$

- Each tree makes up to two splits: GAM with Interactions (GAMI)
- Balances interpretability with interaction modeling
- Most practical choice for interpretable GBDT

# ANOVA Decomposition Process

## Two-Stage Process:

### 1. Aggregation Stage:

- Start with tree ensemble:  $f(x) = \sum_k \eta_k T_k(x)$
- Decompose into leaf nodes:  
$$f(x) = \sum_m v_m \prod_{j \in S_m} I(s_{mj}^l \leq x_j < s_{mj}^u)$$
- Assign effects based on split variables:
  - Main effects: 1 split variable
  - Pairwise interactions: 2 split variables
  - Higher-order: More split variables

### 2. Purification Stage:

- Ensure orthogonality:  $\int f_{i_1 \dots i_t} dx_k = 0$
- Remove identifiability issues
- Center effects to have zero mean
- Cascade from high-order to low-order interactions



# Feature Importance Analysis

## Understanding Overall Feature Impact:

```
1 # Global feature importance
2 result = ts.interpret_fi()
3 result.plot()
```

## Importance Metrics:

- Based on variance of marginal effects
- Normalized to sum to 1
- Accounts for feature scale differences
- Shows relative importance across features

## Interpretation:

- Higher values indicate stronger influence on predictions
- Helps identify key drivers in the model
- Supports feature selection decisions
- Guides business focus areas

# Effect Importance Analysis

## ANOVA Component Analysis:

```
1 # Global effect importance
2 result = ts.interpret_ei()
3 result.plot()
```

## Effect Types:

- **Main effects:** Individual feature contributions
- **Interaction effects:** Pairwise feature interactions
- **Higher-order:** Complex multi-feature interactions

## Applications:

- Understand model complexity
- Identify significant interactions
- Validate domain knowledge
- Guide feature engineering

# Effect Visualization

## Main Effect Plots:

```
1 # Visualize main effect for specific feature
2 result = ts.interpret_effects(features="hr")
3 result.plot()
```

## Plot Interpretation:

- Shows how feature affects predictions
- Reveals non-linear relationships
- Indicates feature importance regions
- Supports business understanding

## Interaction Effect Plots:

```
1 # Visualize pairwise interaction
2 result = ts.interpret_effects(features=("temp", "hum"))
3 result.plot()
```

# Individual Prediction Analysis

## Local Feature Importance:

```
1 # Local interpretation for specific sample
2 result = ts.interpret_local_fi(sample_index=10,
3                               centered=True)
4 result.plot()
```

## Local Effect Importance:

```
1 # ANOVA-based local explanation
2 result = ts.interpret_local_ei(sample_index=10,
3                               centered=True)
4 result.plot()
```

## Local Explanation Components:

- Feature/effect contributions to specific prediction sample
- Comparison to average behavior
- Direction and magnitude of effects

# Introduction to Monotonicity Constraints

## What are Monotonicity Constraints?

Enforce that certain feature-response relationships follow domain knowledge:

### Common Examples:

- **Credit Scoring:** Score should increase with income
- **Risk Assessment:** Risk should decrease with credit rating
- **Real Estate:** Property value should increase with square footage
- **Marketing:** Response rate should increase with previous purchases

**Mathematical Constraint:** For monotonic increasing feature  $x_j$ :

$$\frac{\partial f(x)}{\partial x_j} \geq 0 \quad \forall x \quad (7)$$

**Implementation:** Available in XGBoost through MoDeVa's `MoXGBRegressor` and `MoXGBClassifier`

# Benefits of Monotonicity Constraints

## Interpretability Enhancement:

- **Smoother Functions:** Reduces noise and fluctuations
- **Clearer Patterns:** More interpretable global trends
- **Consistent SHAP:** More stable local explanations
- **Business Logic:** Aligned with domain knowledge

## Model Quality Improvements:

- **Reduced Overfitting:** Constraints act as regularization
- **Better Generalization:** More robust to new data
- **Stable Performance:** Often maintains or improves accuracy

## Interaction with ANOVA:

- Guarantees monotonic shape for constrained features
- Preserves interpretability in decomposition
- Maintains partial monotonicity in interactions
- Simplifies effect visualization

# Model Selection Guidelines

## Choosing the Right GBDT Implementation:

### Use XGBoost when:

- Need monotonicity constraints
- Want extensive hyperparameter control
- Performance is critical
- Working with structured data

### Use LightGBM when:

- Large datasets (100K+ samples)
- Memory is limited
- Training speed matters
- Categorical features present

### Use CatBoost when:

- Many categorical features
- Minimal preprocessing desired
- Overfitting is a concern
- New to GBDT

## Tree Depth Selection:

- **Depth 1:** Maximum interpretability (GAM)
- **Depth 2:** Good balance (GAMI) - *recommended*
- **Depth 3+:** Higher performance, reduced interpretability

# Hyperparameter Tuning Strategy

## Key Parameters to Optimize:

### 1. Tree Structure:

- `max_depth`: Start with 2, increase if needed
- `n_estimators`: Start with 100, use early stopping
- `min_child_weight`: Increase to prevent overfitting

### 2. Learning Control:

- `learning_rate`: Start with 0.1, decrease for more trees
- `subsample`: Try 0.8-1.0 for regularization
- `colsample_bytree`: Use 0.8-1.0 for feature sampling

### 3. Tuning Process:

- ① Fix depth at 2 for interpretability
- ② Tune number of estimators with early stopping
- ③ Optimize learning rate and regularization
- ④ Apply monotonicity constraints if needed



# Key Takeaways

## **MoDeVa Integration:**

- Unified interface across XGBoost, LightGBM, and CatBoost
- Integrated interpretability through Functional ANOVA
- Easy model comparison and deployment

## **Interpretability:**

- Depth constraints enable GAM/GAMI structures
- Global and local explanation methods
- Monotonicity constraints for domain knowledge integration

## **Best Practices:**

- Start with depth-2 trees for interpretability
- Use early stopping and cross-validation
- Apply monotonicity constraints when appropriate
- Validate interpretations with domain experts