

# Model Performance & Residual Analysis

July 5, 2025

# Outline

- 1 Introduction to Performance Evaluation
- 2 Classification Metrics
- 3 Regression Metrics
- 4 Challenges in Performance Measurement
- 5 Introduction to Residual Analysis
- 6 Techniques for Residual Analysis
- 7 Residual Analysis
- 8 Supervised Learning for Residual Analysis
- 9 High-Error Region Identification

# Why Measure Model Performance?

- **Accuracy Assessment:** Quantify how well models predict outcomes for unseen data
- **Model Comparison:** Select the best-performing model among multiple candidates
- **Bias-Variance Analysis:** Detect overfitting (high variance) or underfitting (high bias)
- **Business Impact:** Translate metrics into actionable insights

## Key Concepts

- **Evaluation Metrics:** Classification vs. Regression metrics
- **Data Splitting:** Training, validation, testing datasets
- **Cross-Validation:** Estimate performance across multiple splits
- **Real-World Validation:** Simulate practical scenarios

# Classification Metrics I

## Accuracy

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

Proportion of correctly classified samples

## Recall (Sensitivity)

$$\text{Recall} = \frac{TP}{TP + FN}$$

True positives among all actual positives

## Precision

$$\text{Precision} = \frac{TP}{TP + FP}$$

True positives among predicted positives

## F1-Score

$$F1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

Harmonic mean of precision and recall

# Classification Metrics II

## AUC-ROC

$$\int_0^1 TPR(FPR^{-1}(x))dx$$

Measures model's ability to distinguish between classes

## Log Loss (Cross-Entropy)

$$-\frac{1}{n} \sum_{i=1}^n [y_i \log(p_i) + (1 - y_i) \log(1 - p_i)]$$

Penalizes confident but incorrect predictions

## Brier Score

$$\frac{1}{n} \sum_{i=1}^n (p_i - y_i)^2$$

Evaluates accuracy of predicted probabilities

## When to use each metric:

- Imbalanced classes → Precision, Recall, F1
- Ranking performance → AUC-ROC
- Probability calibration → Log Loss, Brier Score

## Mean Squared Error (MSE)

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

Average squared difference between predictions and actuals

## Mean Absolute Error (MAE)

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |\hat{y}_i - y_i|$$

Average absolute difference between predictions and actuals

## R-Squared (Coefficient of Determination)

$$R^2 = 1 - \frac{\sum_{i=1}^n (\hat{y}_i - y_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

Proportion of variance explained by the model

### When to use each metric:

- Sensitive to outliers → MAE
- Penalize large errors → MSE
- Compare models → R-Squared

# Challenges in Measuring Model Performance

## Data Imbalance

- Metrics like accuracy can be misleading
- Better metrics: precision, recall, F1-score

## Overfitting/Underfitting

- Overfitting: Model performs well on training data but poorly on unseen data
- Underfitting: Insufficient learning of data patterns

## Real-World Applicability

- Lab performance  $\neq$  Real-world performance
- Need for temporal validation or external datasets

## Heteroscedasticity

- Error variance changes across feature values
- May indicate need for transformations

# Performance Evaluation: Model Setup & Training

```
1 # Regression tasks using lightGBM and xgboost
2 from modeva.models import MoLGBMRegressor, MoXGBRegressor
3 # for lightGBM
4 model_lgbm = MoLGBMRegressor(name = "LGBM_model",
5                               max_depth=2,
6                               n_estimators=100)
7 # for xgboost
8 model_xgb = MoXGBRegressor(name = "XGB_model",
9                             max_depth=2,
10                             n_estimators=100)
11
12 # Train models with input: ds.train_x and target: ds.train_y
13 model_lgbm.fit(ds.train_x, ds.train_y)
14 model_xgb.fit(ds.train_x, ds.train_y)
15
```



# Performance Evaluation in ModEva: Reporting

```
1 # Create a testsuite that bundles dataset and model
2 from modeva import TestSuite
3 ts = TestSuite(ds, model_lgbm) # store bundle of dataset and
   model
4
5 # Evaluate performance and summarize into table
6 results = ts.diagnose_accuracy_table(
7     train_dataset="train",
8     test_dataset="test",
9     metric=("MAE", "MSE", "R2")
10 )
11 results.table # Display results in a tabular format
12
```

*This produces a table with metrics for both training and test sets*

# Performance Comparison Between Models

```
1 # Create TestSuite with multiple models
2 tsc = TestSuite(ds, models=[model_lgbm, model_xgb])
3
4 # Performance comparison of 2 models
5 results = tsc.compare_accuracy_table(
6     train_dataset="train",
7     test_dataset="test",
8     metric=("MAE", "MSE", "R2")
9 )
10 results.plot() # Visualize the comparison
11
```

*This generates comparative visualizations to help identify the better performing model*

# What is Residual Analysis?

## Definition

Residual analysis is a diagnostic tool to evaluate model performance by analyzing differences between actual and predicted values.

**For regression:**  $r_i = y_i - \hat{y}_i$

**For classification:** Residuals are expressed as misclassification errors or differences between predicted probabilities and true labels.

## Purpose of Residual Analysis:

- Assess model fit and identify systematic errors
- Examine error distribution patterns
- Detect outliers where model performs poorly
- Identify heteroscedasticity (non-constant variance)
- Validate model assumptions

# Techniques for Residual Analysis

## Residual Plots

- Residuals vs. Predicted:  
Should show random scatter around zero
- Residuals vs. Features:  
Identifies feature-specific weaknesses
- Histogram of Residuals:  
Checks distribution pattern

## Quantile-Quantile (Q-Q) Plot

- Assesses if residuals follow normal distribution
- Compares quantiles of residuals to a theoretical normal distribution

## Interpreting Results

- Random distribution → Model captures data well
- Systematic patterns → Model misspecification
- High variance → Heteroscedasticity
- Large residuals → Outliers or edge cases

## Applications

- Model debugging
- Feature engineering
- Outlier detection
- Assumption validation

# Basic Residual Analysis

```
1 # Create a TestSuite as before
2 ts = TestSuite(ds, model_lgbm)
3
4 # Perform residual analysis with feature "hr"
5 results = ts.diagnose_residual_analysis(
6     features="hr",
7     dataset="train"
8 )
9 results.plot() # Generate residual visualization
10
```

*This produces visualization showing residuals vs. feature "hr"*

- Identify patterns in residuals across different hours of the day
- Check if the model performs consistently throughout the day
- Detect time periods with higher error rates

## Concept

MoDeVa uses supervised ML to analyze residuals, enabling targeted model improvements.

- **Traditional Approach Limitations:**

- Binning and clustering often miss complex non-linear relationships
- Difficult to identify why errors occur in high-dimensional spaces

- **ModEva's Approach:**

- Uses interpretable GBDT models to explicitly model residual errors
- Employs Random Forest proximity matrices for similarity-based clustering
- Identifies high-error regions in feature space
- Provides actionable insights for targeted model improvements

# Methodology 1: Residual Modeling with Interpretable GBDT

- 1 Train interpretable XGBoost model (depth-1 or depth-2) to predict residuals
- 2 Use absolute residuals  $r = |y - \hat{y}|$  as target variable
- 3 Extract feature importance and effects from the residual model

```
1 # Train interpretable GBDT model to analyze residuals
2 results = ts.diagnose_residual_interpret(
3     dataset='test', n_estimators=100, max_depth=2
4 )
5 results.plot("feature_importance") # Plot feature
6                                   importance
7
8 results.plot("effect_importance")  # Plot effect importance
9
10 # Get effect of specific feature on residuals
11 ts_residual = results.value["TestSuite"]
12 ts_residual.interpret_effects("hr", dataset="test").plot()
```

## Methodology 2: Proximity-Based Clustering

- 1 Train Random Forest model on dataset
- 2 Extract proximity matrix measuring similarity between data points
- 3 Cluster similar samples based on this proximity
- 4 Analyze how errors distribute across clusters

```
1 # Cluster based on RF proximity matrix
2 results = ts.diagnose_residual_cluster(
3     dataset="test",
4     response_type="abs_residual",
5     metric="MAE",
6     n_clusters=10,
7     cluster_method="pam",
8     sample_size=2000,
9     rf_n_estimators=100,
10    rf_max_depth=5
11 )
12 results.table # Table of cluster performance
13 results.plot("cluster_residual")
14 results.plot("cluster_performance")
15 results.plot("feature_importance")
16
```



# Identifying & Interpreting High-Error Regions

- 1 Identify clusters with highest average error
- 2 Analyze feature composition of problematic clusters
- 3 Use data drift analysis to compare high-error regions to overall distribution
- 4 Develop targeted interventions to improve model performance

```
1 # Analyze a specific high-error cluster
2 cluster_id = 2 # Cluster with high error rate
3
4 # Compare cluster distribution to overall distribution
5 data_results = ds.data_drift_test(
6     **results.value["clusters"][cluster_id]["data_info"],
7     distance_metric="PSI",
8     psi_method="uniform",
9     psi_bins=10
10 )
11 data_results.plot("summary") # Overall drift summary
12 data_results.plot(name=('density', 'hr')) # Feature
    distribution
13
```

# Actionable Insights from Residual Analysis

## Model Improvements

- **Feature Transformations:** Apply non-linear transformations to features with high residual effect
- **Feature Engineering:** Create new features for high-error regions
- **Model Architecture:** Adjust model complexity or algorithm

## Data Enhancements

- **Targeted Sampling:** Collect more data in high-error regions
- **Outlier Handling:** Develop specific strategies for edge cases
- **Segmented Models:** Create specialized models for challenging subgroups

## Key Takeaway

Residual analysis transforms model evaluation from a simple metric comparison to a targeted diagnostic process that directly informs model improvements.

# Summary: Model Performance & Residual Analysis

- ➊ **Performance Metrics:** Choose appropriate metrics based on problem type and business needs
- ➋ **Residual Analysis:** Critical diagnostic tool for understanding model weaknesses
- ➌ **MoDeVa Framework:** Provides integrated tools for comprehensive model evaluation
- ➍ **Supervised Learning for Residuals:** Powerful approach to identify and interpret error patterns
- ➎ **High-Error Region Identification:** Enables targeted model improvements

## Next Steps

- Apply these techniques to your own models
- Combine performance metrics with residual analysis
- Explore other MoDeVa capabilities for comprehensive model validation