

Hyperparameter Tuning

July 6, 2025

What is Hyperparameter Tuning?

- Hyperparameters are "knobs" that control model training
- Examples: learning rate, max depth, number of estimators
- They are **not learned** from data — must be set manually
- Tuning is the process of finding the best values for these settings

Grid Search: Intuition

- Try all combinations of hyperparameters in a fixed grid
- Like a chef testing every recipe from a list of ingredients
- Works well when search space is small

Objective:

$$\theta^* = \arg \min_{\theta \in \Theta} f(\theta)$$

Algorithm:

- 1 Define grid $\Theta = \theta_1 \times \cdots \times \theta_d$
- 2 For each $\theta \in \Theta$, train model and evaluate $f(\theta)$
- 3 Return best θ^*

Random Search: Intuition

- Sample random combinations of hyperparameters
- Like spinning a roulette wheel to find good combos
- Often outperforms grid search in high dimensions

Objective:

$$\theta^* = \arg \min_{\theta \sim p(\theta)} f(\theta)$$

Algorithm:

- 1 Define sampling distributions p_1, \dots, p_d
- 2 For $i = 1$ to N : sample $\theta^{(i)}$, train model, evaluate
- 3 Return best $\theta^{(i)}$

Particle Swarm Optimization: Intuition

- Inspired by birds or fish searching for food
- Each "particle" (solution) flies in search space
- Particles learn from their own and neighbors' best positions

Velocity Update:

$$v_i(t+1) = wv_i(t) + c_1r_1(p_i - x_i(t)) + c_2r_2(g - x_i(t))$$

Position Update:

$$x_i(t+1) = x_i(t) + v_i(t+1)$$

Where:

- w : inertia weight
- c_1, c_2 : cognitive/social coefficients
- $r_1, r_2 \sim \text{Uniform}(0, 1)$

PSO: Personal Best p and Global Best g

Personal Best p_i :

$$p_i = \arg \min_{x \in \{x_i(1), \dots, x_i(t)\}} f(x)$$

Best position particle i has seen so far (based on validation loss).

Global Best g :

$$g = \arg \min_i f(p_i)$$

Best position found by any particle in the swarm.

Intuition:

- Each particle remembers its own best solution — p_i
- Particles share knowledge to follow the global best — g
- Movement balances exploration (randomness) and exploitation (toward p_i and g)

PSO: Algorithm

- 1 Initialize particles x_i and velocities v_i
- 2 Evaluate $f(x_i)$, update personal best p_i and global best g
- 3 Update velocity and position:
$$v_i \leftarrow \text{update formula}$$
$$x_i \leftarrow x_i + v_i$$
- 4 Repeat for T iterations
- 5 Return best g

Recommended Parameters for PSO

- **Inertia Weight** w : 0.4 – 0.9
 - High w : more exploration
 - Low w : more exploitation
 - Typical value: $w = 0.7$
- **Cognitive Coefficient** c_1 : 1.0 – 2.0 (self-confidence)
- **Social Coefficient** c_2 : 1.0 – 2.0 (swarm confidence)
- **Swarm Size**: 10 – 30 particles
- **Max Iterations**: 20 – 100+
- **Parameter Bounds**: Required for each dimension

Rule of Thumb: Try $w = 0.7$, $c_1 = c_2 = 1.5$, swarm size = 20

Which Method Should You Use?

Criterion	Grid Search	Random Search	PSO
Exploration Strategy	Exhaustive	Random Sampling	Adaptive Swarm Movement
Best For	Small search space	Large/mid space	Complex, continuous space
Search Space	Discrete only	Discrete or continuous	Continuous (preferred)
Computation Cost	High	Moderate	Moderate to High
Convergence	No	No	Yes
Parallelizable	Yes	Yes	Moderate
Requires Bounds	No	Optional	Yes
Smart Feedback	No	No	Yes
Tuning Complexity	Low	Low	Medium
When to Use	Small models or tight budget	High budget, mid-size models	High-dim, noisy objectives

MoDeVa: Hyperparameter Tuning Overview

- MoDeVa supports 3 tuning strategies:
 - `ModelTuneGridSearch`
 - `ModelTuneRandomSearch`
 - `ModelTunePSO`
- Common arguments:
 - `param_distributions`: hyperparameter search space
 - `metric`: e.g., "MSE", "ACC", "F1"
 - `cv`: number of cross-validation folds
 - `n_iter`: only for random search

Example: Grid Search in MoDeVa

```
from modeva.models import ModelTuneGridSearch
param_grid = {"n_estimators": [50, 100, 200],
              "learning_rate": [(i + 1) * 0.01 for i in range(5)]}
model = MoLGBMClassifier(max_depth=2, verbose=-1)
hpo = ModelTuneGridSearch(dataset=ds, model=model)
result = hpo.run(param_grid=param_grid,
                 metric=("AUC", "ACC", "LogLoss", "Brier"),
                 cv=5)

result.table
```

Example: Random Search in MoDeVa

```
from modeva.models import ModelTuneRandomSearch
param_grid = {"alpha": [0.1, 1.0, 10],
              "l1_ratio": [(i + 1) * 0.1 for i in range(10)]}

model = MoElasticNet()
hpo = ModelTuneRandomSearch(dataset=ds, model=model)
result = hpo.run(param_distributions=param_grid,
                 n_iter=20,
                 metric="MSE",
                 cv=5)

result.table
```

Example: PSO Search in MoDeVa

```
from modeva.models import ModelTunePSO
param_bounds = {"max_depth": [1, 4],
                 "learning_rate": [0.01, 1.0]}
param_types = {"max_depth": "int"}

model = MoLGBMClassifier(verbose=-1)
hpo = ModelTunePSO(dataset=ds, model=model)
result = hpo.run(param_bounds=param_bounds,
                  param_types=param_types,
                  n_iter=2,
                  n_particles=10,
                  metric=("AUC", "LogLoss"),
                  cv=5)

result.table
```