

# Reliability Testing

## Evaluation of Prediction Uncertainty

July 5, 2025

# Outline

- 1 Introduction to Model Reliability
- 2 Conformal Prediction Framework
- 3 Reliability Analysis in ModEva
- 4 Identifying Reliability Issues
- 5 Advanced Reliability Analysis
- 6 Remediation Strategies

# What is Model Reliability?

## Definition

Reliability in predictive models refers to their ability to produce consistent and trustworthy outputs by accurately quantifying prediction uncertainty.

## Why Is Reliability Important?

- Enables risk assessment in predictions
- Critical in high-stakes domains like finance, healthcare, or autonomous systems
- Helps users understand when to trust model outputs
- Identifies when a model is "not sure" about its predictions
- Guides data collection and model improvement efforts

## Key Insight

A reliable model is not necessarily more accurate, but it knows when it doesn't know

# Sources of Uncertainty in ML Models

## Data-Related Uncertainty

- Noisy measurements or labels
- Sparse or limited training data
- Missing features or values
- Data quality issues
- Sampling bias

## Model-Related Uncertainty

- Structural limitations
- Parameter uncertainty
- Inherent model limitations
- Optimization challenges

## Environment-Related Uncertainty

- Distribution shift
- Novel or unfamiliar inputs
- Adversarial scenarios
- Edge cases not seen during training

## Predictive Uncertainty Types

- Aleatoric uncertainty (inherent randomness)
- Epistemic uncertainty (model/knowledge limitations)
- Distributional uncertainty (out-of-distribution data)

# Investigating Model Reliability

## Steps for Investigating and Improving Model Reliability:

### ① Uncertainty Quantification

- Use Conformal Prediction to generate prediction intervals
- Provide a range within which true values likely fall
- Establish guaranteed coverage levels

### ② Identification of Less Reliable Regions

- Identify high uncertainty regions in feature space
- Detect areas with outside-of-coverage predictions
- Analyze patterns in unreliable predictions

### ③ Determine Impactful Variables

- Identify key variables contributing to uncertainty
- Quantify feature importance for reliability
- Analyze feature interactions affecting reliability

### ④ Enhance the Model

- Address identified weaknesses
- Apply data-centric and model-centric approaches
- Improve reliability in targeted regions

# Introduction to Conformal Prediction

## Concept

Conformal prediction provides a model-agnostic framework for constructing prediction intervals with guaranteed coverage under minimal assumptions.

## Key Properties:

- Distribution-free coverage guarantees
- Requires only that data are exchangeable (i.i.d. is sufficient)
- Works with any base model or predictor
- For a desired miscoverage rate  $\alpha$ , produces intervals containing the true value with probability  $1 - \alpha$
- Applicable to both classification and regression

## Core Idea

Quantify how different a new test point is from the training data using a nonconformity score

# Full Conformal Prediction

Full Conformal Prediction is the most faithful implementation of the conformal prediction framework. It provides finite-sample, distribution-free validity guarantees under the assumption that the data are exchangeable.

Given a training dataset

$$\mathcal{D} = \{(x_1, y_1), \dots, (x_n, y_n)\}$$

and a new input  $x_{n+1}$ , the goal of full conformal prediction is to construct a prediction set  $\Gamma_n(x_{n+1})$  such that:

$$\mathbb{P}(y_{n+1} \in \Gamma_n(x_{n+1})) \geq 1 - \alpha,$$

where  $\alpha \in (0, 1)$  is a user-defined miscoverage rate.

# Full Conformal Prediction Algorithm (1)

## Step 1: Define a Nonconformity Score

Let  $A(f, (x, y))$  be a nonconformity score function. For each  $i = 1, \dots, n$ , train a model  $f_i$  on  $\mathcal{D}_{-i} = \mathcal{D} \setminus \{(x_i, y_i)\}$ , and compute:

$$\alpha_i = A(f_i, (x_i, y_i)).$$

For example:

- Regression:  $\alpha_i = |f_i(x_i) - y_i|$
- Classification:  $\alpha_i = 1 - \hat{P}_{f_i}(y_i \mid x_i)$

## Step 2: Evaluate Candidate Labels

For a new input  $x_{n+1}$ , and for each candidate label or value  $y \in \mathcal{Y}$ , augment the dataset:

$$\mathcal{D}' = \mathcal{D} \cup \{(x_{n+1}, y)\}.$$

Then, for each  $i = 1, \dots, n + 1$ , train a model on  $\mathcal{D}' \setminus \{(x_i, y_i)\}$ , and compute nonconformity scores:

$$\alpha_i^{(y)} = A(f_i^{(y)}, (x_i, y_i)).$$



# Full Conformal Prediction Algorithm (2)

Let  $\alpha_{n+1}^{(y)}$  denote the nonconformity score for the trial point  $(x_{n+1}, y)$ .

## Step 3: Compute p-value

Define the p-value for the candidate  $y$  as:

$$p(y) = \frac{1}{n+1} \sum_{i=1}^{n+1} \mathbf{1} \left\{ \alpha_i^{(y)} \geq \alpha_{n+1}^{(y)} \right\}.$$

## Step 4: Construct the Prediction Set

The conformal prediction set is:

$$\Gamma_n(x_{n+1}) = \{y \in \mathcal{Y} : p(y) > \alpha\}.$$

Under the assumption of exchangeability of the data:

$$\mathbb{P}(y_{n+1} \in \Gamma_n(x_{n+1})) \geq 1 - \alpha,$$

holds for any finite sample size  $n$ , without distributional assumptions.

# Example of Full Conformal Prediction Calculation (1)

We use logistic regression:

$$\hat{P}(1 \mid x) = \sigma(-2 + 1.5x), \quad \text{where } \sigma(z) = \frac{1}{1 + e^{-z}}$$

Training Data and Nonconformity Scores

$x_i$	$y_i$	$\hat{P}(y_i \mid x_i)$	$\alpha_i = 1 - \hat{P}(y_i \mid x_i)$
0	0	0.881	0.119
1	0	0.623	0.377
2	1	0.731	0.269
3	1	0.924	0.076

Example: Confidence Level 0.8, Test Input  $x = 0.75$

$$z = -2 + 1.5 \cdot 0.75 = -0.875, \quad \sigma(z) \approx 0.294$$

## Example of Full Conformal Prediction Calculation (2)

**Case 1:** Assume  $y = 0$

$$\hat{P}(0 \mid x) = 1 - 0.294 = 0.706, \quad \alpha_5^{(0)} = 1 - 0.706 = 0.294$$

Compare with training scores:  $\{0.119, 0.377, 0.269, 0.076\}$

$$\Rightarrow \# \geq 0.294 = 1 + 1 \text{ (test)} = 2 \Rightarrow p(0) = \frac{2}{5} = 0.4 > \alpha = 0.2$$

**Include  $y = 0$  in prediction set.**

**Case 2:** Assume  $y = 1$

$$\hat{P}(1 \mid x) = 0.294, \quad \alpha_5^{(1)} = 1 - 0.294 = 0.706$$

$$\# \geq 0.706 = 0 + 1 \text{ (test)} = 1 \Rightarrow p(1) = \frac{1}{5} = 0.2 \not> 0.2$$

**Do not include  $y = 1$ .**

$$\Gamma(x = 0.75) = \{0\}$$

# Note on Full Conformal Prediction Algorithm

In the example, for simplicity, we are using approximation where we did not refit the model for each test point and candidate value.

## Full Conformal Prediction

- Requires refitting model for each test point and candidate value
- P-value calculation for each potential output
- Computationally expensive but theoretically optimal
- Strongest coverage guarantees

# Split Conformal Prediction

- Splits data into training and calibration sets
- Train model once on training set
- Compute nonconformity scores on calibration set
- Use empirical quantiles to construct intervals
- Much more computationally efficient
- Slight reduction in statistical efficiency

## Algorithm:

- 1 Split data:  $D_{train}$  and  $D_{cal}$
- 2 Train model on  $D_{train}$
- 3 Compute scores on  $D_{cal}$
- 4 Use score quantiles for intervals

# Conformalized Residual Quantile Regression (CRQR)

## A Powerful Approach for Regression Problems

### 1 Base Model Fitting

- Fit a base model  $f(X)$  to estimate conditional mean
- Calculate residuals:  $r_i = y_i - f(X_i)$

### 2 Residual Quantile Regression

- Train quantile regression models for lower and upper quantiles
- Predict residual quantiles:  $\hat{q}_{\tau_1}(X)$ ,  $\hat{q}_{\tau_2}(X)$
- Capture heteroscedasticity (changing variance across feature space)

### 3 Conformalization

- Compute nonconformity scores on calibration set
- Adjust quantile predictions to ensure coverage guarantee

### 4 Prediction Intervals

- For new point  $X$ , construct interval:
- $C(X) = [f(X) + \hat{q}_{\tau_1}(X) - Q_{1-\alpha}(s), f(X) + \hat{q}_{\tau_2}(X) + Q_{1-\alpha}(s)]$
- Combines flexibility of quantile regression with conformal guarantees

# Nonconformity Scores

## For Regression:

- Absolute residual:  $s_i = |y_i - \hat{f}(X_i)|$
- Normalized residual:  $s_i = \frac{|y_i - \hat{f}(X_i)|}{\hat{\sigma}(X_i)}$
- CRQR score:  $s_i = \max\{q_{\tau_1}(X_i) - r_i, r_i - q_{\tau_2}(X_i)\}$

## For Binary Classification:

- Score based on predicted probabilities:
- $s_i = \begin{cases} 1 - \hat{p}_i & \text{if } y_i = 1 \\ \hat{p}_i & \text{if } y_i = 0 \end{cases}$
- Lower score indicates better conformity with training data

## Advantages and Considerations:

- Distribution-free coverage guarantees
- Works with any base predictor
- Split conformal trades statistical efficiency for computational efficiency
- Can capture heteroscedasticity with appropriate scoring functions
- Can be extended to handle covariate shift with weighted approaches

# Basic Reliability Assessment in ModEva

```
1 # Create a testsuite that bundles dataset and model
2 from modeva import TestSuite
3 ts = TestSuite(ds, model_lgbm) # store bundle of dataset and
   model
4
5 # reliability assessment using Split Conformal Prediction
6 results = ts.diagnose_reliability(
7     train_dataset="test", test_dataset="test",
8     test_size=0.5, alpha=0.1,
9     max_depth=5)          # depth for quantile regression
   model
10 results.table
11
```

*This generates a table showing prediction interval width and coverage*

## Key Configuration

**alpha=0.1** means we expect 90% of true values to fall within prediction intervals



# Understanding Reliability Issues

## Two Key Aspects of Conformal Prediction Outputs:

### 1. High Uncertainty

- Measured by prediction interval width
- For regression: wide intervals indicate high uncertainty
- Flag points with width threshold:  $W(X) > Q_{1-\beta}(W)$
- $\beta$  typically 0.9 for top 10% widest intervals
- For classification: empty or multi-class prediction sets

### 2. Coverage Violations

- Occurs when true value falls outside prediction interval
- For regression:  $y \notin C_{n,\alpha}(X)$
- For classification: true class not in prediction set
- Expected violation rate should match  $\alpha$
- Higher-than-expected violations indicate model issues

## Common Patterns Indicating Model Weakness:

- High uncertainty + good coverage: Model is honest about uncertainty
- High uncertainty + poor coverage: Model may be misspecified
- Low uncertainty + poor coverage: Model is overconfident

# Analyzing Reliability Issues in MoDeVa

```
1 # Analyze distribution differences between less reliable
   regions and overall data
2 data_results = ds.data_drift_test(
3     **results.value["data_info"],
4     distance_metric="PSI", psi_method="uniform", psi_bins
       =10)
5
6 # Display summary of distribution differences ranked by PSI
7 data_results.plot("summary")
8 # Visualize density comparison for a specific feature
9 data_results.plot(("density", "hr"))
10
```

*This ompares the distribution of features in unreliable regions vs. the overall dataset*

## Key outputs:

- PSI summary: Features ranked by distribution difference
- Density plots: Visual comparison of distributions
- Histogram comparisons: Bin-level differences

# Feature Slicing for Reliability Analysis

```
1 # Univariate slicing: analyze reliability across a single
   feature
2 results = ts.diagnose_slicing_reliability(
3     features="hr",           # feature to analyze
4     train_dataset="train", test_dataset="test",
5     test_size=0.5, metric="width")
6 results.plot()
7
```

*This shows how prediction interval width varies across different hours of the day*

## What to look for:

- Features or feature values with unusually high uncertainty
- Consistent patterns in uncertainty across feature ranges
- Times of day, categories, or value ranges with reliability issues

# Multiple Feature Analysis

```
1 # Analyze multiple features independently
2 results = ts.diagnose_slicing_reliability(
3     features=((("hr",), ("atemp",), ("weekday",))),
4     train_dataset="train", test_dataset="test",
5     test_size=0.5, metric="coverage")
6 # View results for each feature
7 results.plot("hr")
8 results.plot("atemp")
9 results.plot("weekday")
10
```

*This analyzes coverage across multiple features independently*

## Coverage Analysis

Comparing actual coverage to expected coverage ( $1-\alpha$ ) reveals regions where the model is systematically over- or under-confident

# Feature Interaction Analysis

```
1 # Two-dimensional slicing: analyze feature interactions
2 results = ts.diagnose_slicing_reliability(
3     features=("hr", "atemp"),      # feature pair to analyze
4     train_dataset="train",
5     test_dataset="test",
6     test_size=0.5,
7     random_state=0)
8 results.plot()
9
```

*This visualizes how combinations of feature values affect reliability*

## Benefits:

- Identifies complex interactions affecting model reliability
- Reveals conditional dependencies in uncertainty
- Shows where feature combinations create problematic regions
- Helps detect subtle reliability patterns missed in univariate analysis

# Model Comparison

```
1 # Compare reliability between models
2 tsc = TestSuite(ds, models=[model_lgbm, model_xgb])
3
4 # Compare overall reliability metrics
5 results = tsc.compare_reliability(
6     train_dataset="train", test_dataset="test",
7     test_size=0.5, alpha=0.1, max_depth=5)
8 results.table
9
```

*This compares reliability metrics between different models*

## What to compare:

- Average prediction interval width (narrower is better, if coverage is maintained)
- Actual coverage (closer to target  $1-\alpha$  is better)
- Consistency of coverage across feature space
- Trade-off between interval width and coverage

# Supervised Machine Learning for Uncertainty Analysis

```
1 # Use Random Forest clustering with prediction interval
   width as target
2 results = ts.diagnose_residual_cluster(
3     dataset="test", response_type="pi_width", metric="MAE",
4     n_clusters=10, cluster_method="pam", sample_size=2000,
5     rf_n_estimators=100, rf_max_depth=5) # RF parameters
6 results.table
7 results.plot()
8
```

*This uses Random Forest proximity to cluster similar samples based on prediction uncertainty*

## Key outputs:

- Cluster table: Performance metrics for each cluster
- Feature importance: Variables driving uncertainty clusters
- Cluster visualization: Similarity patterns in high-uncertainty regions

# Detailed Cluster Analysis

```
1 # Analyze a specific high-uncertainty cluster
2 cluster_id = 2 # cluster with high uncertainty
3
4 # Compare cluster distribution to overall distribution
5 data_results = ds.data_drift_test(
6     **results.value["clusters"][cluster_id]["data_info"],
7     distance_metric="PSI", psi_method="uniform", psi_bins
8     =10)
9 data_results.plot("summary")
10 data_results.plot(name=('density', 'atemp'))
```

*This analyzes the feature distribution patterns of a specific high-uncertainty cluster*

## Insights from cluster analysis:

- Distinct feature patterns in high-uncertainty regions
- Feature interaction effects not visible in univariate analysis
- Natural groupings of similar uncertainty patterns
- Key drivers of prediction uncertainty clusters



# Remediation: Data-Centric Approaches

## 1. Targeted Data Augmentation

- Focus on high-uncertainty regions
- Collect additional samples in weak regions
- Use active learning to select informative samples
- Prioritize areas with low coverage

## 2. Feature Engineering

- Create interaction terms for regions with nonlinear patterns
- Develop domain-specific features for high-uncertainty areas
- Transform features to better capture heteroscedasticity
- Add features that help discriminate in uncertain regions

### Key Principle

Targeted data improvements in high-uncertainty regions can significantly enhance model reliability

# Remediation: Model-Centric Approaches

## 1. Local Model Enhancement

- Train specialized models for unreliable regions
- Implement segment-specific models
- Use Mixture of Experts (MoE) approach
- Weight models based on local performance

## 2. Architecture Modifications

- Add capacity in high-uncertainty regions
- Try alternative modeling frameworks

## 3. Loss Function Adjustments

- Weight samples from uncertain regions higher
- Implement reliability-focused penalties
- Balance overall performance with local improvements

## 4. Ensemble Strategies

- Combine models with complementary reliability profiles
- Weight ensemble components based on local uncertainty
- Implement model switching based on detected uncertainty

# Implementation Framework

## Diagnose

- Apply conformal prediction
- Identify high-uncertainty regions
- Detect coverage violations
- Analyze feature patterns

## Prioritize

- Focus on most unreliable regions
- Rank features by importance
- Consider business impact
- Balance effort vs. improvement

## Implement & Validate

- Apply targeted remediation
- Rerun reliability testing
- Compare before/after metrics
- Iterate as needed

## Systematic Approach

Improving reliability requires understanding uncertainty patterns, applying targeted interventions, and validating improvements

# Summary: Model Reliability Testing

- **Understanding Reliability:** A model's ability to produce consistent outputs with appropriate uncertainty estimates
- **Conformal Prediction:** Framework for creating prediction intervals with guaranteed coverage
- **Reliability Analysis:** Identifying regions with high uncertainty or coverage violations
- **Feature Analysis:** Using slicing and clustering to understand patterns in unreliable predictions
- **Model Comparison:** Different models may show varying reliability profiles
- **Targeted Remediation:** Combining data-centric and model-centric approaches to improve reliability

## Key Takeaway

Reliable models provide appropriate uncertainty estimates, allowing users to make informed decisions about when to trust the model's predictions